

1. Describe and analyze an algorithm to figure out what Mr. Fox says.

Solution: Let $Chickens(i, r, d)$ denote the maximum number of chickens that Mr. Fox can earn starting at the i th booth, assuming he just said “Ring!” r times in a row, or he just said “Ding!” d times in a row. Notice that we must have either $r = 0$ or $d = 0$, and Boggis, Bunce, and Bean’s rules imply that $r \leq 3$ and $d \leq 3$. We need to compute $Chickens(1, 0, 0)$.

The $Chickens$ function satisfies the following recurrence:

$$Chickens(i, r, d) = \begin{cases} 0 & \text{if } i > n \\ Chickens(i + 1, 1, 0) + A[i] & \text{if } i \leq n \text{ and } d = 3 \\ Chickens(i + 1, 0, 1) - A[i] & \text{if } i \leq n \text{ and } r = 3 \\ \max \begin{cases} Chickens(i + 1, r + 1, 0) + A[i] \\ Chickens(i + 1, 0, d + 1) - A[i] \end{cases} & \text{otherwise} \end{cases}$$

We can memoize this function into a three-dimensional array $Chickens[1..n, 0..3, 0..3]$. Because each entry $Chickens[i, r, d]$ in this array depends only on entries of the form $Chickens[i + 1, \cdot, \cdot]$, we can fill this array in order of decreasing i , filling in the 6 interesting entries $Chickens[i, r, d]$ for each i in arbitrary order. The resulting algorithm runs in $O(n)$ time. ■

Rubric: Standard dynamic programming rubric. This is not the only correct linear-time solution.

2. (a) Describe an algorithm to compute the length of the shortest palindrome supersequence of a given string.

Solution: Let $A[1..n]$ be the input string. For any indices i and j , let $SPS(i, j)$ denote the length of the shortest palindrome supersequence of the substring $A[i..j]$. This function obeys the following recurrence:

$$SPS(i, j) = \begin{cases} 0 & \text{if } i > j \\ 1 & \text{if } i = j \\ 2 + \min \begin{cases} SPS(i+1, j) \\ SPS(i, j-1) \end{cases} & \text{if } i < j \text{ and } A[i] \neq A[j] \\ 2 + \min \begin{cases} SPS(i+1, j-1) \\ SPS(i+1, j) \\ SPS(i, j-1) \end{cases} & \text{otherwise} \end{cases}$$

We need to compute $SPS(1, n)$.

We can memoize the function LPS into a two-dimensional array. Each entry depends on the $LPS[i, j]$ depends on (at most) three entries $LPS[i+1, j]$, $LPS[i, j-1]$, and $LPS[i+1, j-1]$ immediately below and/or to the left. Thus, we can fill the array from bottom to top in the outer loop, and from left to right in inner loop, as follows:

```
SPS(A[1..n]):
  for i ← n down to 1
    SPS[i, i-1] ← 0
    SPS[i, i] ← 1
    for j ← i+1 to n
      SPS[i, j] ← min {2 + SPS[i+1, j], 2 + SPS[i, j-1]}
      if A[i] = A[j]
        SPS[i, j] ← min {SPS[i, j], 2 + SPS[i+1, j-1]}
  return SPS[1, n]
```

The resulting algorithm runs in $O(n^2)$ time. ■

Rubric: 5 points: standard dynamic programming rubric (scaled). This is not the only correct evaluation order. This is not the only correct algorithm with this running time.

- (b) Describe an algorithm to compute the length of the shortest dromedrome supersequence of a given string.

Solution: We actually solve a more general problem. A *shortest common supersequence* of two strings w and x is a string z of smallest possible length such that w is a subsequence of z and x is a subsequence of z . For example, **ITRONMYSTANRK** is a shortest common supersequence of **IRONMAN** and **TONYSTARK**. Every dromedrome supersequence of $A[1..n]$ is a string of the form ww , where w is a common supersequence of some prefix $A[1..j-1]$ and the corresponding suffix $A[j..n]$.

For any indices i, j, k , let $SCS(i, j, k)$ denote the length of the shortest common supersequence of the prefix $A[1..i]$ and the substring $A[j..k]$. The length of the shortest dromedrome supersequence of A is exactly $2 \cdot \max_{1 \leq j \leq n} SCS(j-1, j, n)$. The SCS function obeys the following recurrence:

$$SCS(i, j, k) = \begin{cases} i & \text{if } k < j \\ k - j + 1 & \text{if } i < 1 \\ \min \begin{cases} 1 + SCS(i, j, k-1) \\ 1 + SCS(i-1, j, k) \end{cases} & \text{if } A[i] \neq A[k] \\ \min \begin{cases} 1 + SCS(i, j, k-1) \\ 1 + SCS(i-1, j, k) \\ 1 + SCS(i-1, j, k-1) \end{cases} & \text{if } A[i] = A[k] \end{cases}$$

We can memoize this function into a 3-dimensional array $SCS[0..n, 0..n, 0..m]$. Each entry $SCS[i, j, k]$ depends only on entries $SCS[i', j, k']$ where $i' < i$ or $k' < k$ (or both); in particular, the second index j is unchanged. Thus, we can fill the array with three nested for-loops, considering j in arbitrary order in the outer loop, increasing i in the middle loop, and increasing k in the inner loop. We only need to consider entries $SCS[i, j, k]$ where $0 \leq i \leq j-1 \leq k \leq n$. The resulting algorithm runs in $O(n^3)$ time.

MINDROME DROME($A[1..n]$):

$mindds \leftarrow 2n$ *⟨easy upper bound⟩*

for $j \leftarrow 1$ to n

for $k \leftarrow j-1$ to n

$SCS[0, j, k] \leftarrow 0$

for $i \leftarrow 1$ to $j-1$

$SCS[i, j, j-1] \leftarrow 0$

for $k \leftarrow j$ to n

$SCS[i, j, k] \leftarrow \min\{1 + SCS[i, j, k-1], 1 + SCS[i-1, j, k]\}$

if $A[i] = A[k]$

$SCS[i, j, k] \leftarrow \min\{SCS[i, j, k], 1 + SCS[i-1, j, k-1]\}$

$mindds \leftarrow \min\{mindds, 2 \cdot SCS[j-1, j, n]\}$

return $mindds$

■

Rubric: 5 points: standard dynamic programming rubric (scaled). This is not the only correct evaluation order. We can simplify this algorithm very slightly by using a two-dimensional memoization array, indexed only by i and k .

3. Suppose you are given a DFA M with k states for Jeff's favorite regular language $L \subseteq (\emptyset+1)^*$.
- (a) Describe and analyze an algorithm that decides whether a given bit-string belongs to the language L^* .

Solution (8/10): Use the text-segmentation algorithm in the textbook, implementing $\text{IsWORD}(i, j)$ by tracing the substring $w[i..j]$ through the DFA in constant time per character. The segmentation algorithm makes $O(n^2)$ calls to IsWORD , and each call to IsWORD runs in $O(n)$ time, so the overall algorithm runs in $O(n^3)$ time. ■

Solution (10/10): We use the text-segmentation algorithm in the textbook, implementing the function IsWORD using the following dynamic programming algorithm. Assume that the DFA M is represented by a pair of arrays $\delta[1..k, \emptyset..1]$ and $A[1..k]$, as shown on page 2 of the DFA lecture notes, and that the start state s is state 1. Let $w[1..n]$ denote the input string.

For any indices i and j , let $\text{state}(i, j) = \delta^*(s, w[i..j])$; more explicitly, $\text{state}(i, j)$ is the state that the substring $w[i..j]$ reaches in our given DFA. This function obeys the following recurrence (directly mirroring the recursive definition of δ^*):

$$\text{state}(i, j) = \begin{cases} 1 & \text{if } j < i \\ \delta[\text{state}(i, j-1), w[j]] & \text{otherwise} \end{cases}$$

We can memoize this function into a two-dimensional array $\text{state}[1..n, 0..n]$. Because each entry $\text{state}[i, j]$ depends only on the previous entry $\text{state}[i, j-1]$ in the same row, we can fill this array in standard row-major order in $O(n^2)$ time.

After we've precomputed the state array, we can implement the helper function IsWORD as follows:

```
IsWORD(i, j):
    return A[state[i, j]]
```

The text segmentation makes $O(n^2)$ calls to IsWORD , and IsWORD runs in $O(1)$ time, so the overall algorithm runs in $O(n^2)$ time. ■

- (b) Describe and analyze an algorithm that partitions a given bit-string into as many substrings as possible, such that L contains every substring in the partition. Your algorithm should return only the number of substrings, not their actual positions.

Solution: We modify the text segmentation algorithm in the textbook as follows.

For any index i , let $MinWords(i)$ denote the minimum size of any partition of the suffix $w[i..n]$ into words, where “words” are defined by the boolean function $IsWord$. This function obeys the following recurrence:

$$MinWords(i) = \begin{cases} 0 & \text{if } i > n \\ \min \{1 + MinWords(j+1) \mid IsWord(i, j)\} & \text{otherwise} \end{cases}$$

(To handle the case where the suffix $w[i..n]$ cannot be split into words, we define $\min \emptyset = \infty$.) Like the function *Splittable*, this function can be memoized into a one-dimensional array $MinWords[1..n+1]$, which we can fill in *decreasing* index order. The resulting algorithm makes $O(n^2)$ calls to $IsWord$.

```

MINWORDS( $w[1..n]$ ):
   $MinWords[n+1] \leftarrow 0$ 
  for  $i \leftarrow n$  down to 1
     $MinWords[i] \leftarrow \infty$ 
    for  $j \leftarrow i$  to  $n$ 
      if  $IsWord(i, j)$ 
         $MinWords[i] \leftarrow \min \left\{ \begin{array}{l} MinWords[i] \\ 1 + MinWords[j+1] \end{array} \right\}$ 
  return  $MinWords[1]$ 

```

Now we proceed exactly as in part (a). If we implement each call to $IsWord$ by simulating the DFA, the resulting algorithm runs in $O(n^3)$ time. However, if we precompute all values of the *state* function, then each call to $IsWord$ can be evaluated in $O(1)$ time, and thus the overall algorithm runs in $O(n^2)$ *time*. ■